

Capstone Project- Credit Card Default- Taiwanese Company

Uzair Faruqi

2020-11-24

Credit Card Default Data- for a Taiwanese Company

In This Project I have Worked with the Taiwanese Credit card Data set. This data set has 30000 rows and 24 columns. The objective is to estimate the probability of default payment by credit card client using the data provided. These attributes are related to various details about a customer, his past payment information and bill statements. The Data set is available on the github repository <https://github.com/uzairfaru2/Credit-Card-Default> it is also available on the UCI Machine Learning repository.

I am anyways picking up the data from the repository and hence the code should run directly

Synopsis

The Objective of the assignment is to predict whether a particular customer will default payment next month or not. The result is an extremely valuable piece of information for the bank to take decisions regarding offering credit to its customer and could affect the bank's revenue.

The data set is a tricky one as it has a mix of categorical and continuous variables also it is a classic case of an imbalanced data set and hence any model developed on the entire data set would have an element of bias as defined by the overall Accuracy which these models which many of the Machine learning Models will try and Maximize.

I would be trying to address this issue using various means. First I will try and develop a normal logistic regression model to predict the probability default, measuring the probability at two different benchmarks of 0.5 and 0.25 and see the overall accuracy as well as the specificity and sensitivity.

Then next I would be developing a decision tree model using the rpart library to achieve the objective. In the decision tree I would be utilizing these 4 methods available in the rpart package to address the issue 1. Under Sampling the training set 2. Changing the prior probability 3. Assigning a Cost to Default 4. Changing the weights to default vs non Default. As we see that when we develop a model on the entire set, the Accuracy might seem good but the specificity or Sensitivity gets compromised.

With the Various versions of the respective Tree models as we see that we can balance the Objectives in terms of selecting the models which has a Better Sensitivity or Specificity.

We will round it off by developing a Boosting model i.e a Gradient Boosting model, Multi Variate Adaptive Recursive Spline - Better Known as xgboost and Adaboost and see the confusion matrix developed on the same. The Boosting models have all been trained on the Undersampled Training set where we have made attempts to balance the data set. I have also tried to develop the model on the full data set. I would be comparing the performance of the various models or the various approaches applied.

As we see that we would have to strike a balance between the Specificity and the Sensitivity of the models to achieve the Objectives.

Loading the necessary libraries and Extracting the Data

```
library(tidyverse)
library(broom)
library(caret)
library(ipred)
library(lattice)
library(ggplot2)
library(caret)
library(MASS)
library(car)
library(knitr)
library(mice)
library(lattice)
library(reshape2)
library(skimr)
library(ROSE)
library(rpart)
library(rpart.plot)
library(earth)
```

```
data <- read.csv('https://raw.githubusercontent.com/uzaifaru2/Credit-Card-Default/main/credit_card_defa
```

```
# Exploring the data
```

```
dim(data)
```

```
[1] 30000 25
```

```
# We are having 30,000 records with around 25 columns - variables ( including a customer ID column)
# The defaults have been encoded as 1 and 0 otherwise
```

```
# All the variables including the customer ID seems to have been considered as integers
```

```
# Looking at the data to see if any Inconsistency
```

```
table(data$EDUCATION)
```

```
0      1      2      3      4      5      6
```

```
14 10585 14030 4917 123 280 51
```

```
# There are data with 5 and 6 defined for Education which hasn't been defined in the data
# data understanding document. We might have to group them under 4 which is others
```

```
table(data$MARRIAGE)
```

```
0      1      2      3
```

54 13659 15964 323

here also we have many under 0 , which has not been defined , we can put it under 3 Which is Others

```
#replace 0's with NAN, replace others too
data$EDUCATION[data$EDUCATION == 0] <- 4
data$EDUCATION[data$EDUCATION == 5] <- 4
data$EDUCATION[data$EDUCATION == 6] <- 4
data$MARRIAGE[data$MARRIAGE == 0] <- 3

# Checking on whether changes have been Made
```

```
table(data$EDUCATION)
```

1	2	3	4
---	---	---	---

10585 14030 4917 468

```
table(data$MARRIAGE)
```

1	2	3
---	---	---

13659 15964 377

About the Data

There are around 30,000 data points i.e around 30,000 unique data points All the data including the customer ID have been considered as integers We would have to make the necessary corrections, converting some of the data into factors and some of the variables might have to be one hot encoded so that we can run the models relate to the Boosting Algorithms

Creating Derived Variables

One of the Important Tasks of a Data scientist is to understand the data and to work on making certain derived variables which would be more effective as variables then the Variables themselves

For Example the Billing Amount or the credit balance might not be so useful themselves say as compared to say the Ratio of Billing Amount to the Credit balance .

There could be many such variables that can be derived and used in the model developed to improve on its accuracy. Much of it depends on the ingenuity of the Data Scientists and the understanding of the domain by them.

Also while we might derive many variables there would have to be certain data engineering to be done to address anomalies etc

```
# Let us work out certain Derived variables
# One can be as innovative as possible in Deriving these variables and these variables
# Can actually effect the performance of the model itself also.

# The Derived Variables that we can look at are
```

```

# The increase in the Billing amount in september as compared to August
# The Billing amount in a percentage of the Credit limit
# Payment to Bill AMNT Ration

data2 <-data %>% mutate( inc_billing_amnt = (BILL_AMT1- BILL_AMT2)/BILL_AMT2 ,
                        bill_amnt_Limit = BILL_AMT1/LIMIT_BAL,
                        pay_bill_amnt_ratio = PAY_AMT1/BILL_AMT2)

# Now Another Problem , what needs to be done for cases which are Giving NA's or Inf/-inf

# One Option is to remove these variables another option might be to replace the NA's by
# 0's or the maximum values
# Infy can be handled by replacing the divisor ( if it was a 0 by the median value) but it will end
# up changing a variable itself which we should avoid.
# Depending on the prevalence of these values

# write.csv(data2, file = "data2.csv")

# We are going to replace the NA's by 0 and the Infinities by the maximum value in that column

max(data2$inc_billing_amnt[is.finite(data2$inc_billing_amnt)])

```

```
[1] 13009.38
```

```

data2$inc_billing_amnt[data2$inc_billing_amnt == Inf] <- max(data2$inc_billing_amnt
[is.finite(data2$inc_billing_amnt)])
data2$inc_billing_amnt[data2$inc_billing_amnt == -Inf] <- min(data2$inc_billing_amnt
[is.finite(data2$inc_billing_amnt)])
data2$inc_billing_amnt[is.na(data2$inc_billing_amnt) ] <- 0

# managing the Infinity Values for the pay_bill_amnt_ratio both the negative and positive Infinities

max(data2$pay_bill_amnt_ratio[is.finite(data2$pay_bill_amnt_ratio)])

```

```
[1] 4444.333
```

```
min(data2$pay_bill_amnt_ratio[is.finite(data2$pay_bill_amnt_ratio)])
```

```
[1] -497.8
```

```

data2$pay_bill_amnt_ratio[data2$pay_bill_amnt_ratio == Inf] <- max(data2$pay_bill_amnt_ratio[is.finite(
data2$pay_bill_amnt_ratio[data2$pay_bill_amnt_ratio == -Inf] <- min(data2$pay_bill_amnt_ratio[is.finite(
data2$pay_bill_amnt_ratio[is.na(data2$pay_bill_amnt_ratio) ] <- 0

# Checking for NA's

sapply(data2, function(x) sum(is.na(x)))

```

ID	LIMIT_BAL	SEX	EDUCATION
0	0	0	0
MARRIAGE	AGE	PAY_0	PAY_2
0	0	0	0
PAY_3	PAY_4	PAY_5	PAY_6
0	0	0	0
BILL_AMT1	BILL_AMT2	BILL_AMT3	BILL_AMT4
0	0	0	0
BILL_AMT5	BILL_AMT6	PAY_AMT1	PAY_AMT2
0	0	0	0
PAY_AMT3	PAY_AMT4	PAY_AMT5	PAY_AMT6
0	0	0	0
default_payment	inc_billing_amnt	bill_amnt_Limit	pay_bill_amnt_ratio
0	0	0	0

Data Engineering and creating the Training set and the Test Set

Now we would be working on making the data ready for the models being developed. This would involve handling the NA's if any, Doing One Hot Encoding since I intend to use the Boosting algorithms to measure their efficacy. Handling the Numerical data as well as factors since though the target variable is stored as a 1 or a 0 it is stored as integer which we might have to convert into factors for the Algo to work properly.

```
# lets remove the ID columns from the data set
data2 <- data2[ , -1]

# We see that certain variables which should be categorical are stored as integers or
# numeric manner
# These should be converted into factors for further analysis
# The variables which should be factors are " Sex", " Education " , " Marriage",
# Using them as numeric would give a false indication of their relevance for the End
# Many of the variables like the History of the payment status can also be used as numeric
# but as numeric they are not significant as conveyed by the correlation matrix and
# hence we are using # these are factors

names <- c(2:4)

data2[,names] <- lapply(data2[,names] , factor)
# str(data2)
#One-Hot Encoding
# Creating dummy variables is converting a categorical variable to as many binary variables as
# where are categories.
dmy <- dummyVars(" ~ .", data = data2, fullRank = T)

train_transformed <- data.frame(predict(dmy, newdata = data2))
# See the structure of the new dataset
# str(train_transformed)

#Converting the dependent variable back to categorical
train_transformed$default_payment<-as.factor(train_transformed$default_payment)
```

```
# checking the same and the other variables
#str(train_transformed)

#Splitting training set into two parts based on outcome: 75% and 25%
set.seed(7)
index <- createDataPartition(train_transformed$default_payment, p = 0.75, list=FALSE, times = 1)
trainSet <- train_transformed[ index,]
testSet <- train_transformed[-index,]
```

Now we have the Training set Ready and the Test Set, we will start with developing a Simple Logistics regression model and see the Results and we can also observe the Implications of an Imbalanced data

Logistic Regression Model

Logistic regression is the First model usually attempted in any Classification problem. It is the Linear regression equivalent model as far as classification problem is concerned.

In short, Logistic Regression is used when the dependent variable(target) is categorical. For example:

To predict whether an email is spam (1) or not spam (0) Whether the tumor is malignant (1) or not (0)

It is named as ‘Logistic Regression’, because it’s underlying technique is quite the same as Linear Regression. There are structural differences in how linear and logistic regression operate. Therefore, linear regression isn’t suitable to be used for classification problems.

Its name is derived from one of the core function behind its implementation called the logistic function or the sigmoid function. It’s an S-shaped curve that can take any real-valued number and map it into a value between 0 and 1, but never exactly at those limits.

```
log.model <- glm(default_payment ~., data = trainSet, family = binomial)

summary(log.model)
```

Call: glm(formula = default_payment ~ ., family = binomial, data = trainSet)

Deviance Residuals: Min 1Q Median 3Q Max
-3.1122 -0.6900 -0.5502 -0.2768 3.8256

Coefficients: Estimate Std. Error z value Pr(>|z|)

(Intercept) -8.034e-01 1.073e-01 -7.484 7.19e-14 ***LIMIT_BAL -1.500e-06 2.299e-07 -6.524 6.85e-11***
SEX.2 -1.361e-01 3.561e-02 -3.822 0.000133 EDUCATION.2 -8.311e-02 4.121e-02 -2.017 0.043705
EDUCATION.3 -1.356e-01 5.544e-02 -2.447 0.014425 *
EDUCATION.4 -1.115e+00 2.173e-01 -5.131 2.88e-07 MARRIAGE.2 -1.785e-01 4.030e-02 -4.429
9.48e-06 MARRIAGE.3 -1.278e-01 1.498e-01 -0.853 0.393689
AGE 5.434e-03 2.173e-03 2.501 0.012398 *
PAY_0 5.907e-01 2.063e-02 28.628 < 2e-16 PAY_2 8.221e-02 2.384e-02 3.448 0.000566 PAY_3
1.004e-01 2.660e-02 3.773 0.000161 PAY_4 1.356e-02 2.889e-02 0.469 0.638798
PAY_5 5.860e-02 3.122e-02 1.877 0.060540 .
PAY_6 -8.698e-03 2.584e-02 -0.337 0.736407
BILL_AMT1 -4.647e-06 1.371e-06 -3.390 0.000699 BILL_AMT2 3.967e-06 1.682e-06 2.358
0.018353 *
BILL_AMT3 7.636e-08 1.513e-06 0.050 0.959737
BILL_AMT4 1.131e-06 1.499e-06 0.755 0.450428
BILL_AMT5 -7.089e-07 1.722e-06 -0.412 0.680613
BILL_AMT6 7.665e-07 1.378e-06 0.556 0.578029

```
PAY_AMT1 -1.425e-05 2.624e-06 -5.431 5.61e-08 PAY_AMT2 -7.568e-06 2.338e-06 -3.237
0.001208 PAY_AMT3 -3.483e-06 1.998e-06 -1.744 0.081241 .
PAY_AMT4 -2.245e-06 1.906e-06 -1.178 0.238892
PAY_AMT5 -2.088e-06 1.898e-06 -1.100 0.271249
PAY_AMT6 -1.794e-06 1.480e-06 -1.213 0.225255
inc_billing_amnt 2.621e-05 8.337e-06 3.143 0.001671 bill_amnt_Limit -3.550e-01 7.354e-02
-4.828 1.38e-06 ** pay_bill_amnt_ratio -9.398e-05 1.400e-04 -0.671 0.502046
— Signif. codes: 0 ‘ ’ 0.001 ‘ ’ 0.01 ‘ ’ 0.05 ‘ ’ 0.1 ‘ ’ 1
```

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 23779 on 22499 degrees of freedom

Residual deviance: 20765 on 22470 degrees of freedom AIC: 20825

Number of Fisher Scoring iterations: 5

Summarising From The Model

By Creating the Model and Looking at the Summary of the Model we can see which of the variables are important , i.e the ones with the stars

- Bill Amount to the Limit Ratio
- INcrease in Billing Amount
- Payment Amount 2 - Amount Paid 2 months back
- Payment Amount 1 - Payment made 1 Month Back
- Bill Amount last Month
- Bill Amount 2 months back
- Marriage
- Education
- Increase in Billing Amount as compared to previous month
- Increase in Billing Amount as a Ratio

These variables are important as far as creating and running a model is concerned.

I had run the model on these set of Important variables as well as the Model with all the variables . There wasnt any noticeable difference in the output and hence sticking with all the variables for this project.

```
# Two of the derived variables etc

log.predictions <- predict(log.model, testSet, type="response")

log.predictions.rd <- ifelse(log.predictions > 0.5, 1, 0)
log.predictions.rd <- as.factor(log.predictions.rd)

con_matrix_log_0.5 <- confusionMatrix(log.predictions.rd, testSet$default_payment)

# With This Type of Mode The Sensitivity Seems to be High but the Precision seems to be lower.
# Accuracy is around 80 % but that could be the case , but we could have got an accuracy of around
# 75 % by predicting non Defaults
# Let us see how we can change things , by Changing the default probability from 0.5 to say 0.25
```

```

log.predictions.rd1 <- ifelse(log.predictions > 0.25, 1, 0)
log.predictions.rd1 <- as.factor(log.predictions.rd1)

con_matrix_log_0.25 <- confusionMatrix(log.predictions.rd1, testSet$default_payment)

# Extracting the Accuracy , Sensitivity and Specificity from the Confusion Matrix Created

log.accuracy.0.5 <- con_matrix_log_0.5$overall[["Accuracy"]]
log.accuracy.0.25 <- con_matrix_log_0.25$overall[["Accuracy"]]
log.specificity.0.5 <- con_matrix_log_0.5$byClass[["Specificity"]]
log.specificity.0.25 <- con_matrix_log_0.25$byClass[["Specificity"]]
log.sensitivity.0.5 <- con_matrix_log_0.5$byClass[["Sensitivity"]]
log.sensitivity.0.25 <- con_matrix_log_0.25$byClass[["Sensitivity"]]

```

The Output of the Logistic Regression

The Accuracy obtained is around 80.6 % and the Specificity 23.6 % for the cut off @ 0.4

The Accuracy obtained is around 75.6 % and the Specificity increases to 54.8 % for the cut off @ 0.4

Here for our case the the Positive is non default and the negative is default (Not Referring to the 1's and 0s')

Hence the specificity measures our efficiency in terms of identifying out of the total number of defaulters how well we have been able to identify them.

We would be evaluation various measures to do that specifically using various measures as given below

Tree Based Models

Now we would be working with some Tree based Models to see how they perform

Sensitivity and specificity both increases depending on the cut off chosen , accuracy Decreases, it might all depend on the priority of the organization

These are the typical characteristics of an imbalanced data set

We will now go through a set of measures designed to address these type of Imbalanced data

These Measures are -Undersampling (Since the Data set is sufficiently large we can as of now avoid Over sampling) -The Undersampling would be done only on the Training Set -Changing the Prior Probabilities -Including a Loss Matrix

These are some of the methods that can be utilised to address this problem of imbalanced data there are many other methods also.

Creating an Undersample Data set from the Training Set We would be using the Rose Library of R to do this and to develop the respective Machine learning model.

First Let us Look at the Proportions of the Default in the Training Set

```
table(trainSet$default_payment)
```

```
0      1
```

```
17523 4977
```

```
prop.table(table(trainSet$default_payment))
```

```
0      1
```

```
0.7788 0.2212
```

we See that the Proportion of defaults is around 22 % , which is a Unbalanced Dataset we will try and Create an UNder sample Data set where the Proportion of defaults and the Proportion of Non Defaults are equal and we would be training our models on this Balanced Data set and the testing would be done on the Test Data set Developed

```
train_under_sample <- ovun.sample(default_payment ~ ., data = trainSet,  
                                   method = "under", N = 9954, seed = 1)$data  
table(train_under_sample$default_payment)
```

```
0 1 4977 4977
```

Here we are building the models using the undersample data set

```
# Now Lets start Building Tree Based Models and observe whether we are able to improve on  
# The Accuracy or the other measures like Specificity
```

```
tree_undersample <- rpart(default_payment ~ ., method = "class",  
                           data = train_under_sample, control = rpart.control(cp = 0.001))
```

```
# Now Let us Plot the Tree's
```

```
#plot(tree_undersample, uniform = TRUE)
```

```
# Add labels to the decision tree  
# text(tree_undersample)
```

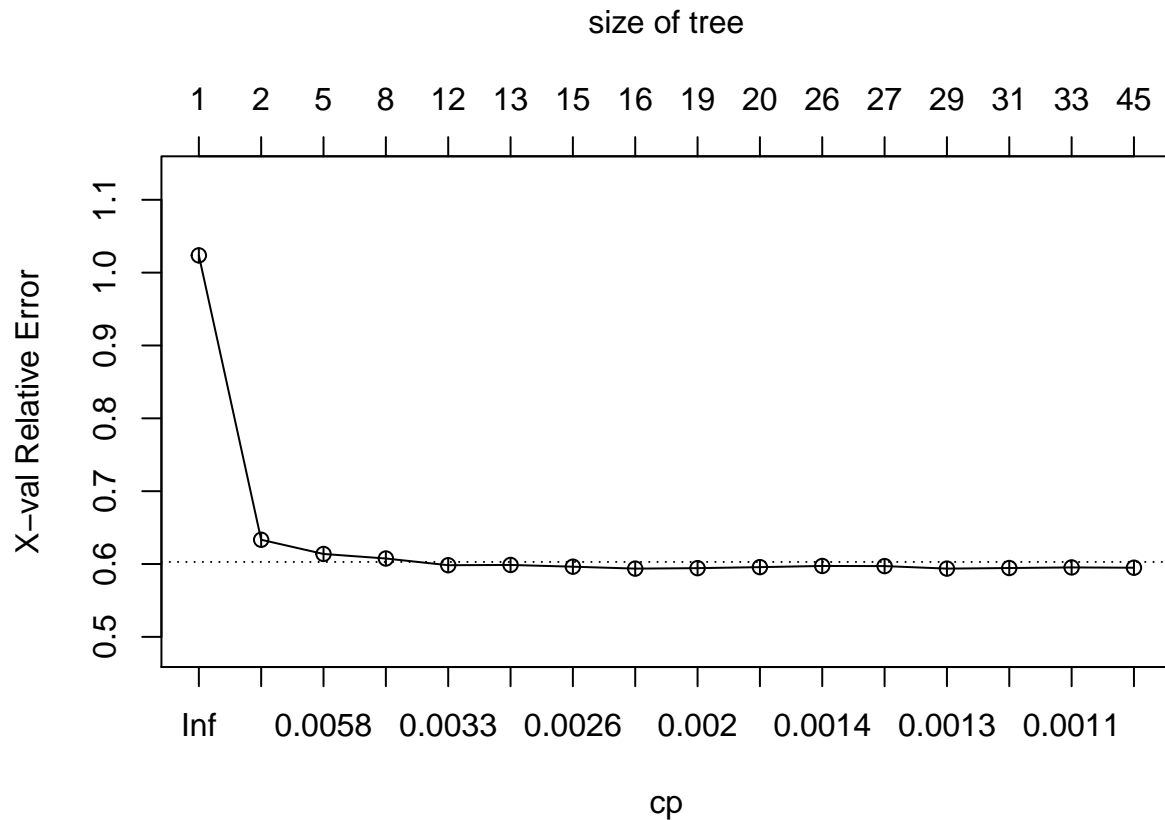
```
# I am avoiding PLOtting the Trees , since the Image isn't coming out  
# properly in the RMD generated File ,  
# it Can be seen properly in a normal R Studio Console
```

Pruning the Trees

We are now pruning the trees to have least error and to avoid over fitting

```
# Now Lets Prune the tree for the Model tree_undersample

# Plotting the cross-validated error rate as a function of the complexity parameter
plotcp(tree_undersample)
```



```
# Using printcp() to identify for which complexity parameter
# the cross-validated error rate is minimized.

printcp(tree_undersample)
```

Classification tree: rpart(formula = default_payment ~ ., data = train_under_sample, method = "class", control = rpart.control(cp = 0.001))

Variables actually used in tree construction: [1] AGE bill_amnt_Limit BILL_AMT1 BILL_AMT2
 [5] BILL_AMT3 BILL_AMT4 BILL_AMT5 BILL_AMT6
 [9] EDUCATION.2 inc_billing_amnt LIMIT_BAL PAY_0
 [13] PAY_2 PAY_3 PAY_4 PAY_6
 [17] PAY_AMT1 PAY_AMT2 PAY_AMT3 PAY_AMT4
 [21] PAY_AMT5 PAY_AMT6

Root node error: 4977/9954 = 0.5

n= 9954

CP	nsplit	rel error	xerror	xstd
----	--------	-----------	--------	------

```

1 0.3666868 0 1.00000 1.02371 0.0100203 2 0.0076351 1 0.63331 0.63331 0.0093249 3 0.0043534 4 0.61041
0.61382 0.0092455 4 0.0038176 7 0.59735 0.60759 0.0092192 5 0.0028129 11 0.57685 0.59835 0.0091791 6
0.0027125 12 0.57404 0.59875 0.0091808 7 0.0024111 14 0.56862 0.59634 0.0091702 8 0.0021097 15 0.56620
0.59373 0.0091586 9 0.0018083 18 0.55957 0.59433 0.0091613 10 0.0014467 19 0.55777 0.59574 0.0091676 11
0.0014065 25 0.54692 0.59735 0.0091747 12 0.0013060 26 0.54551 0.59715 0.0091738 13 0.0012055 28 0.54290
0.59373 0.0091586 14 0.0011051 30 0.54049 0.59453 0.0091622 15 0.0010046 32 0.53828 0.59534 0.0091658 16
0.0010000 44 0.52381 0.59494 0.0091640

```

```

# Creating an index for of the row with the minimum xerror
index0 <- which.min(tree_undersample$cptable[ , "xerror"])

# Creating tree_min
tree_min <- tree_undersample$cptable[index0, "CP"]

# Prune the tree using tree_min
ptree_undersample <- prune(tree_undersample, cp = tree_min)

```

Changing the Prior Probabilities

in the rpart package the prior probabilities by default are taken as per their proportions , we can sort of trick R by modifying the same in the parms function to say 60/40

```

tree_prior <- rpart(default_payment ~ ., method = "class",parms = list(prior=c(0.6, 0.4)),
                    control = rpart.control(cp = 0.001),
                    data = trainSet)

# Plot the decision tree
# plot(tree_prior, uniform = TRUE)

# Add labels to the decision tree
# text(tree_prior)

```

Pruning the Tree

Pruning the Tree for the Model with the Changed Prio Probabilities

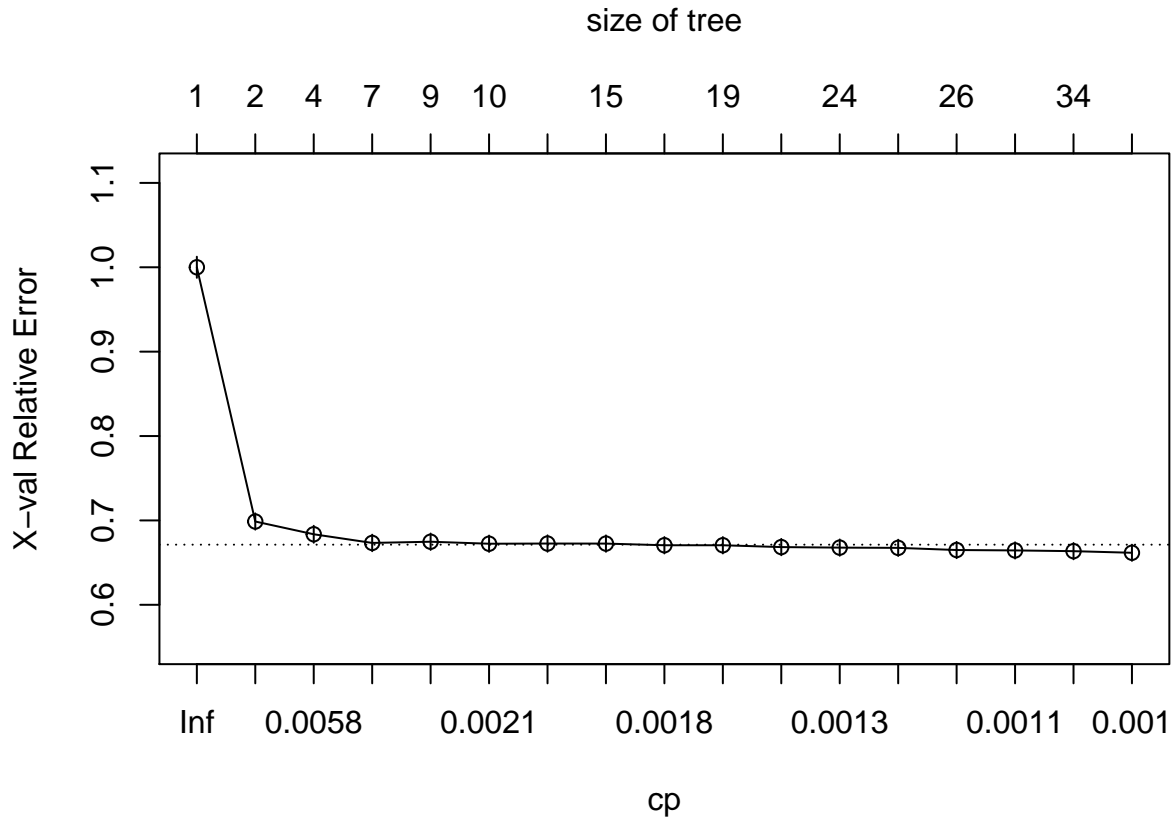
```

# As the Trees get more complex, they are not clearly visible( The Nodes)

# Pruning the Trees with Changed Prior Probabilities

    # Plot the cross-validated error rate as a function of the complexity parameter
plotcp(tree_prior)

```



```
# Use printcp() to identify for which complexity parameter the
# cross-validated error rate is minimized.
printcp(tree_prior)
```

Classification tree: `rpart(formula = default_payment ~ ., data = trainSet, method = "class", parms = list(prior = c(0.6, 0.4)), control = rpart.control(cp = 0.001))`

Variables actually used in tree construction: [1] AGE BILL_AMT1 BILL_AMT4
 [4] inc_billing_amnt LIMIT_BAL MARRIAGE.2
 [7] PAY_0 PAY_2 PAY_3
 [10] PAY_4 PAY_5 PAY_AMT1
 [13] PAY_AMT2 PAY_AMT3 PAY_AMT4
 [16] PAY_AMT5 pay_bill_amnt_ratio

Root node error: $9000/22500 = 0.4$

n= 22500

	CP	nsplit	rel error	xerror	xstd
1	0.3012985	0	1.00000	1.00000	0.0125092
2	0.0100265	1	0.69870	0.69870	0.0096775
3	0.0033079	3	0.67865	0.67865	0.0096775
4	0.0025746	6	0.66872	0.67330	0.0098142
5	0.0022597	8	0.66358	0.67474	0.0097367
6	0.0019117	9	0.66132	0.67237	0.0097038
7	0.0018915	12	0.65545	0.67259	0.0097267
8	0.0018500	14	0.65166	0.67253	0.0097278
9	0.0017414	15	0.64981	0.67056	0.0097216
10	0.0016745	18	0.64459	0.67056	0.0097216
11	0.0013292	22	0.63789	0.66838	0.0097104
12	0.0012840	23	0.63656	0.66777	0.0096641
13	0.0011984	24	0.63528	0.66748	0.0096589
14	0.0011105	25	0.63408	0.66491	0.0096518
15	0.0011060	29	0.62906	0.66442	0.0096523
16	0.0010082	33	0.62448	0.66357	0.0096511
17	0.0010000	34	0.62347	0.66166	0.0095991

```

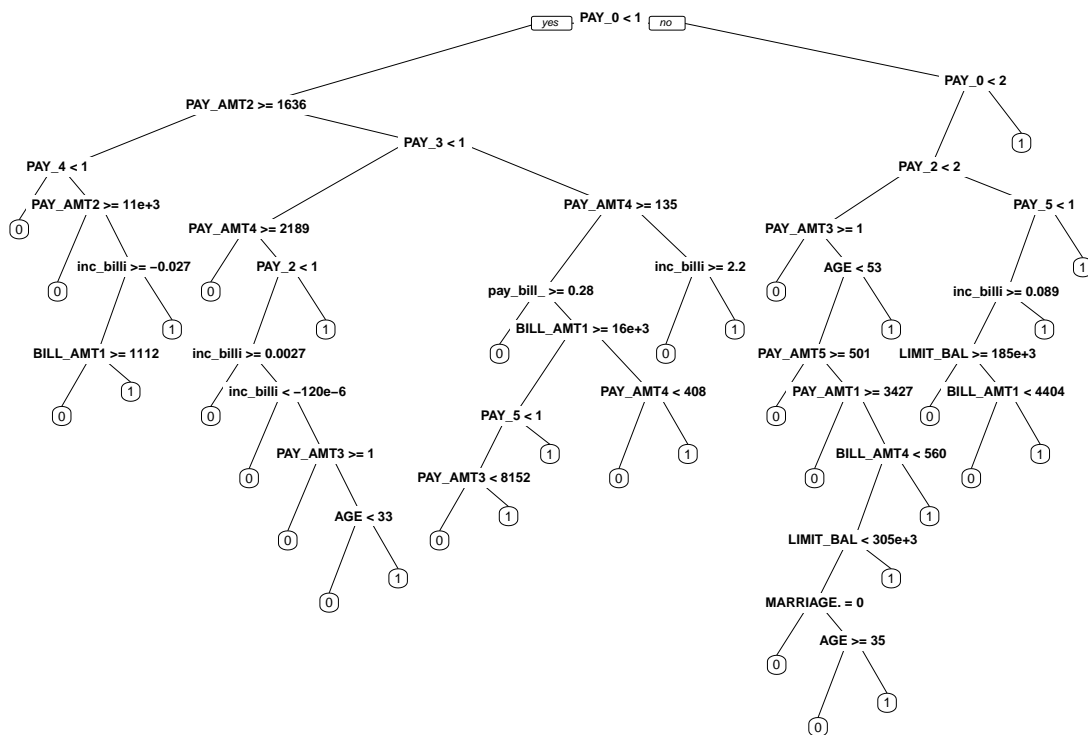
# Create an index for of the row with the minimum xerror
index <- which.min(tree_prior$cptable[ , "xerror"])

# Create tree_min
tree_min <- tree_prior$cptable[index, "CP"]

# Prune the tree using tree_min
ptree_prior <- prune(tree_prior, cp = tree_min)

# Use prp() to plot
prp(ptree_prior)

```



Tree with a Loss Matrix Included

We are now developing a Tree by including a loss matrix where we are penalizing a Default 10 Times more as compared to Non Default. An Arbitrary Number chosen. The Company can choose any other number depicting the actual loss that they might face in case of a default.

```

tree_loss_matrix <- rpart(default_payment ~ ., method = "class",
  parms = list(loss = matrix(c(0, 10, 1, 0), ncol=2)),
  control = rpart.control(cp = 0.001), data = trainSet)

# Plot the decision tree

```

```
#plot(tree_loss_matrix, uniform = TRUE)

# Add labels to the decision tree
#text(tree_loss_matrix)
```

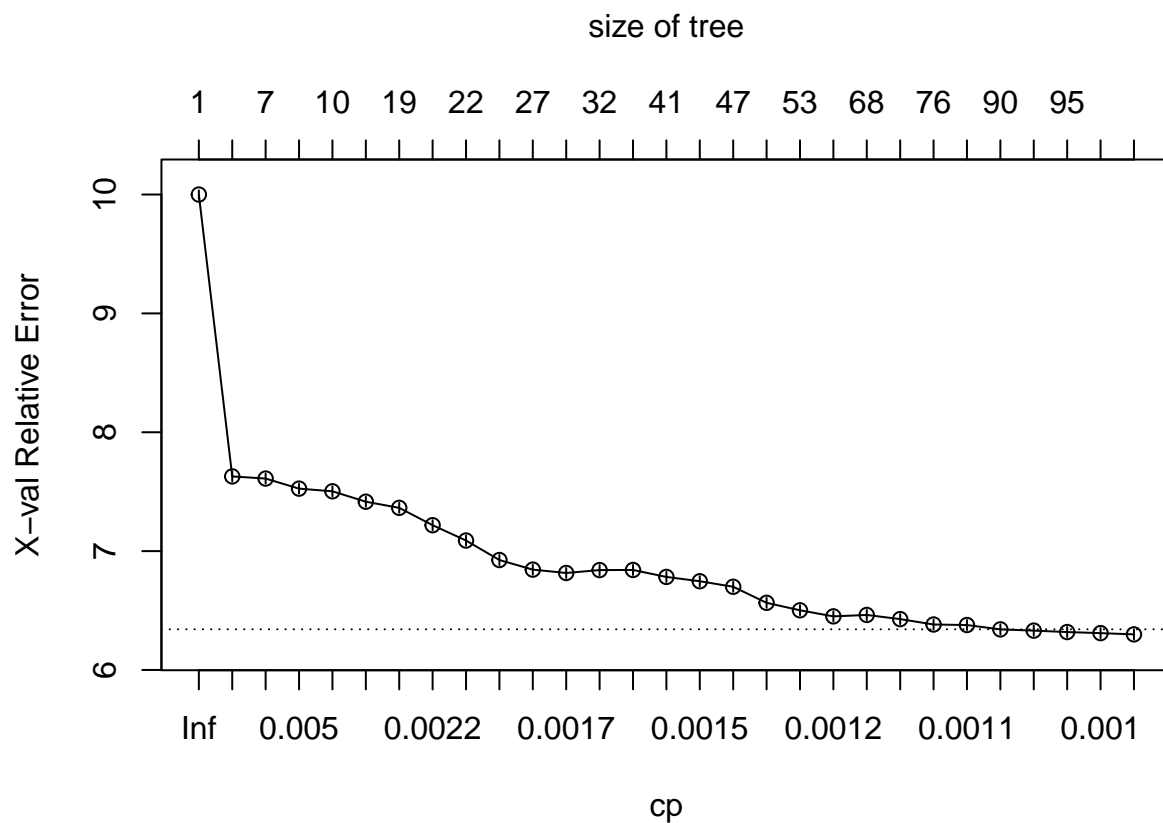
Pruning the Tree

Pruning the Tree for the Model developed by including a Loss Matrix

```
# Pruning the tree with the loss matrix

set.seed(345)
tree_loss_matrix <- rpart(default_payment ~ ., method = "class", data = trainSet,
                           parms = list(loss=matrix(c(0, 10, 1, 0), ncol = 2)),
                           control = rpart.control(cp = 0.001))

# Plot the cross-validated error rate as a function of the complexity parameter
plotcp(tree_loss_matrix)
```



```
printcp(tree_loss_matrix)
```

Classification tree: rpart(formula = default_payment ~ ., data = trainSet, method = "class", parms = list(loss = matrix(c(0, 10, 1, 0), ncol = 2)), control = rpart.control(cp = 0.001))

Variables actually used in tree construction: [1] AGE bill_amnt_Limit BILL_AMT1
 [4] BILL_AMT2 BILL_AMT3 BILL_AMT4
 [7] BILL_AMT5 BILL_AMT6 EDUCATION.2
 [10] inc_billing_amnt LIMIT_BAL PAY_0
 [13] PAY_3 PAY_4 PAY_5
 [16] PAY_6 PAY_AMT1 PAY_AMT2
 [19] PAY_AMT3 PAY_AMT4 PAY_AMT5
 [22] PAY_AMT6 pay_bill_amnt_ratio SEX.2

Root node error: $17523/22500 = 0.7788$

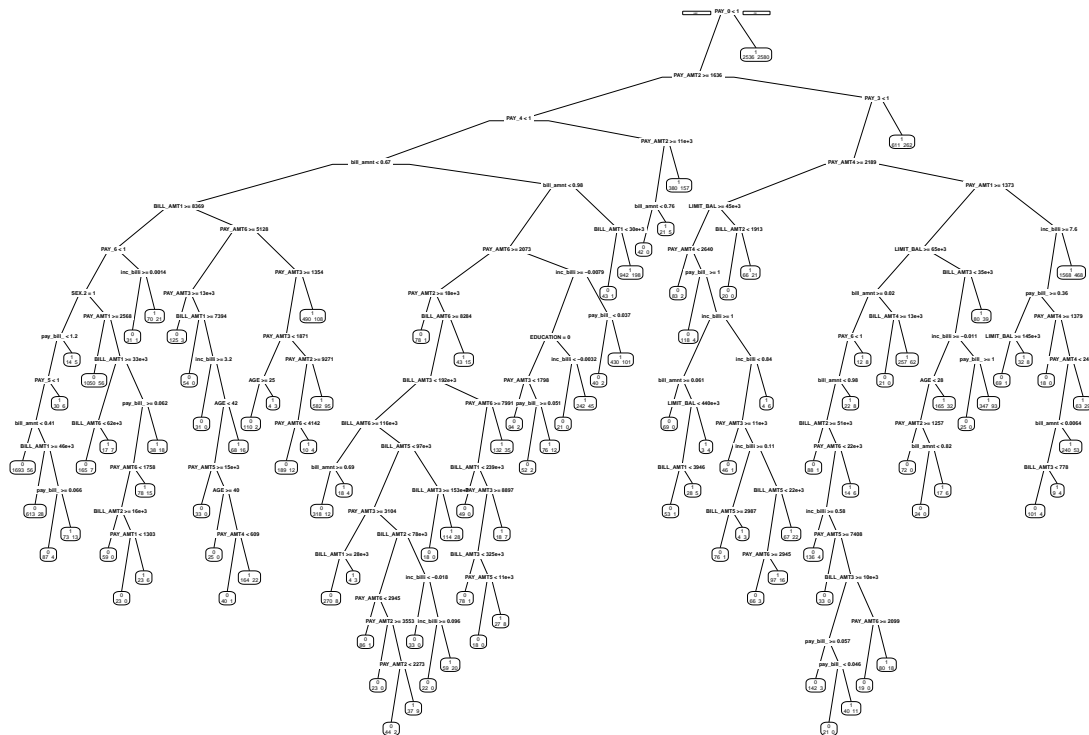
n= 22500

	CP	nsplit	rel error	xerror	xstd
1	0.0128260	0	1.00000	10.0000	0.035529
2	0.0068481	5	0.90675	7.6284	0.041937
3	0.0067911	6	0.89990	7.6104	0.041962
4	0.0036238	7	0.89311	7.5257	0.042063
5	0.0027849	9	0.88586	7.5031	0.042088
6	0.0027107	16	0.86349	7.4156	0.042181
7	0.0024539	18	0.85807	7.3639	0.042230
8	0.0020544	20	0.85316	7.2181	0.042357
9	0.0018832	21	0.85111	7.0891	0.042447
10	0.0018262	23	0.84734	6.9251	0.042531
11	0.0017691	26	0.84147	6.8441	0.042563
12	0.0017120	28	0.83793	6.8160	0.042573
13	0.0016550	31	0.83279	6.8404	0.042563
14	0.0016093	34	0.82783	6.8413	0.042561
15	0.0014838	40	0.81676	6.7827	0.042577
16	0.0014457	42	0.81379	6.7463	0.042586
17	0.0013126	46	0.80797	6.6999	0.042598
18	0.0012840	50	0.80272	6.5649	0.042619
19	0.0012555	52	0.80015	6.5020	0.042624
20	0.0012127	56	0.79513	6.4509	0.042622
21	0.0011984	67	0.77692	6.4628	0.042623
22	0.0011414	73	0.76973	6.4278	0.042621
23	0.0011223	75	0.76745	6.3824	0.042615
24	0.0010843	85	0.75621	6.3781	0.042613
25	0.0010653	89	0.75187	6.3415	0.042608
26	0.0010558	92	0.74867	6.3309	0.042606
27	0.0010272	94	0.74656	6.3190	0.042604
28	0.0010082	98	0.74245	6.3099	0.042603
29	0.0010000	107	0.73252	6.2995	0.042599

```
# Create an index for of the row with the minimum xerror
index2 <- which.min(tree_loss_matrix$cptable[ , "xerror"])

# Create tree_min
tree_min2 <- tree_loss_matrix$cptable[index2, "CP"]
# Prune the tree using cp = 0.0012788
ptree_loss_matrix <- prune(tree_loss_matrix, cp = tree_min2)

# Use prp() and argument extra = 1 to plot the pruned tree
prp(ptree_loss_matrix, extra = 1)
```



Tree with Weights

Here we are developing Tree model by Assigning different weights to default vs Non Default

```
# Pruning the Tree for the case with Weights
# This vector contains weights of 1 for the non-defaults in the training set,
# and weights of 3 for defaults in the training sets.
# By specifying higher weights for default, the model will assign higher
# importance to classifying defaults correctly.
```

```
# Creating a Vector of weights
```

```
case_weights <- ifelse(trainSet$default_payment == 1, 3, 1 )
```

```
head(case_weights)
```

```
[1] 3 3 1 1 1 1
```

```
str(case_weights)
```

```
num [1:22500] 3 3 1 1 1 1 1 1 1 1 ...
```

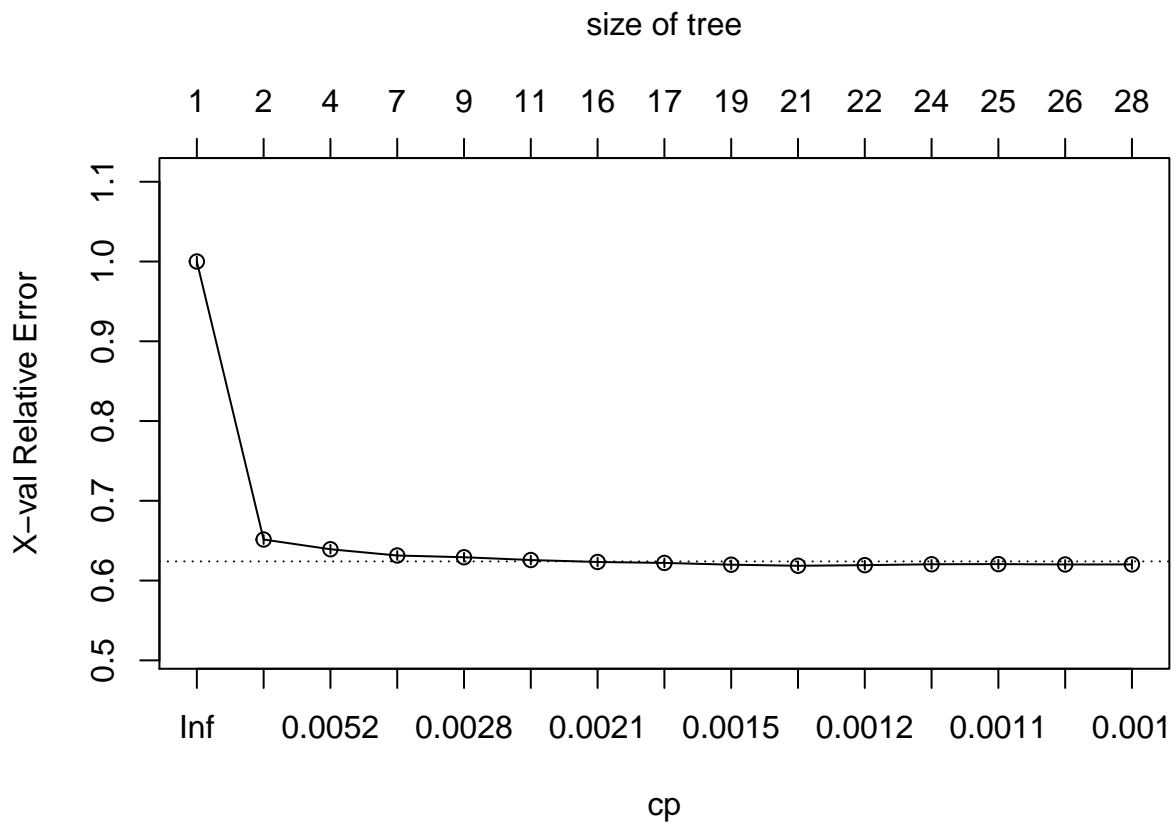


```
head(trainSet$default_payment)
```

```
[1] 1 1 0 0 0 0 Levels: 0 1
```

```
set.seed(345)
tree_weights <- rpart(default_payment ~ ., method = "class",
                      data = trainSet, weights = case_weights,
                      control = rpart.control(minsplit = 5, minbucket = 2, cp = 0.001))

# Plot the cross-validated error rate for a changing cp
plotcp(tree_weights)
```

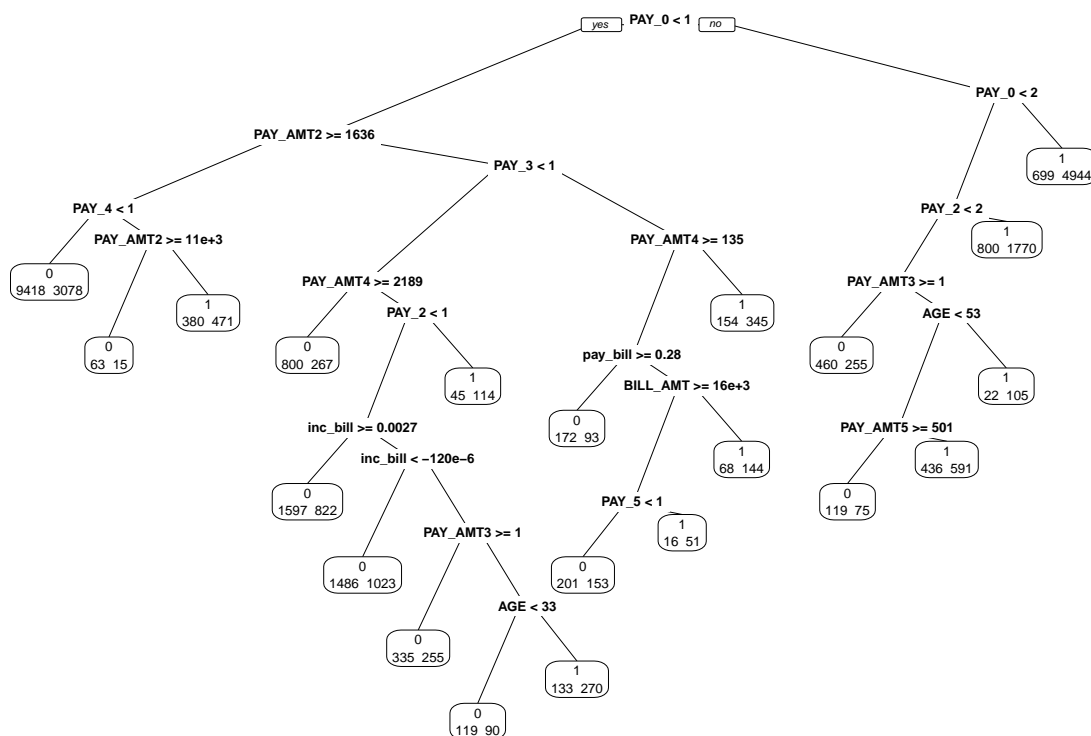


```
# Create an index for of the row with the minimum xerror
index <- which.min(tree_weights$cp[, "xerror"])

# Create tree_min
tree_min <- tree_weights$cp[index, "CP"]

# Prune the tree using tree_min
ptree_weights <- prune(tree_weights, tree_min)

# Plot the pruned tree using the rpart.plot()-package
prp(ptree_weights, extra = 1)
```



Testing of the Different Tree Based Models

Now we would be predicting the probabilities of defaults on the different models by using the Predict Function and by Printing out the Confusion Matrixes

*# Now Lets Predict on the test set using the different models that we have developed and
#Compute the Confusion Matrixes for the respective models*

Make predictions for each of the pruned trees using the test set.

```
pred_undersample <- predict(ptree_undersample, newdata = testSet, type = "class")
pred_prior <- predict(ptree_prior, newdata = testSet, type = "class")
pred_loss_matrix <- predict(ptree_loss_matrix, newdata = testSet, type = "class")
pred_weights <- predict(ptree_weights, newdata = testSet, type = "class")
```

*# Now Lets Create the Confusion Matrixes of the Resulting Models and observe the Accuracies
#and the Specificity/Sensitivity Obtained*

```
pred_undersample <- as.factor(pred_undersample)
pred_prior <- as.factor(pred_prior)
pred_loss_matrix <- as.factor(pred_loss_matrix)
pred_weights <- as.factor(pred_weights)
```

```
con_matrix_undersample <- confusionMatrix(pred_undersample, testSet$default_payment)
con_matrix_prior <- confusionMatrix(pred_prior, testSet$default_payment)
```

```

con_matrix_loss_matrix <- confusionMatrix(pred_loss_matrix, testSet$default_payment)
con_matrix_weights <- confusionMatrix(pred_weights, testSet$default_payment)

# Extracting the Accuracy, Specificity and the Sensitivity

ptree_undersample.accuracy <- con_matrix_undersample$overall[["Accuracy"]]
ptree_undersample.specificity <- con_matrix_undersample$byClass[["Specificity"]]
ptree_undersample.sensitivity <- con_matrix_undersample$byClass[["Sensitivity"]]

ptree_prior.accuracy <- con_matrix_prior$overall[["Accuracy"]]
ptree_prior.specificity <- con_matrix_prior$byClass[["Specificity"]]
ptree_prior.sensitivity <- con_matrix_prior$byClass[["Sensitivity"]]

ptree_loss_matrix.accuracy <- con_matrix_loss_matrix$overall[["Accuracy"]]
ptree_loss_matrix.specificity <- con_matrix_loss_matrix$byClass[["Specificity"]]
ptree_loss_matrix.sensitivity <- con_matrix_loss_matrix$byClass[["Sensitivity"]]

ptree_weights.accuracy <- con_matrix_weights$overall[["Accuracy"]]
ptree_weights.specificity <- con_matrix_weights$byClass[["Specificity"]]
ptree_weights.sensitivity <- con_matrix_weights$byClass[["Sensitivity"]]

## All the Accuracies, Specificity and Sensitivity has been stored in
# separate variables since I intend to compile these and print out a summarized sheet

```

Boosting Algorithms

Unlike many ML models which focus on high quality prediction done by a single model, boosting algorithms seek to improve the prediction power by training a sequence of weak models, each compensating the weaknesses of its predecessors.

boosting is a generic algorithm rather than a specific model. Boosting needs you to specify a weak model (e.g. regression, shallow decision trees, etc) and then improves it.

We would be using two Boosting Algorithms

MARS

Multivariate Adaptive Regression Splines

Multivariate Adaptive Regression Splines, or MARS, is an algorithm for complex non-linear regression problems.

The algorithm involves finding a set of simple linear functions that in aggregate result in the best predictive performance. In this way, MARS is a type of ensemble of simple linear functions and can achieve good performance on challenging regression problems with many input variables and complex non-linear relationships.

AdaBoost

AdaBoost is a specific Boosting algorithm developed for classification problems (also called discrete AdaBoost). The weakness is identified by the weak estimator's error rate.

In each iteration, AdaBoost identifies miss-classified data points, increasing their weights (and decrease the weights of correct points, in a sense) so that the next classifier will pay extra attention to get them right.

Creating the Training Set and the Test Set for the Boosting Algorithms

We would be creating the models for both the undersample data as well as the Full data set

```
# Creating the Training Set and the Test Set
# Here we will be scaling the Data to see the performance of the models with Scaled Data
# We Try and Build These models on the Entire set as well as the Undersampled data set
# and Observe The accuracies and the Specificity/Sensitivity Achieved

# Training the Complete Data set

trainSet1 <- trainSet %>% mutate(default_payment =
                                factor(ifelse(default_payment == "1", "Yes", "No")))
preProcess_range_model <- preProcess(trainSet1, method='range')
trainSet1 <- predict(preProcess_range_model, newdata = trainSet1)

train_under_sample1 <- train_under_sample %>% mutate(default_payment =
                                                       factor(ifelse(default_payment == "1", "Yes", "No")))

preProcess_range_model1 <- preProcess(train_under_sample1, method = 'range')
train_under_sample1 <- predict(preProcess_range_model1, newdata = train_under_sample1)

# Preparing the test Set

testSet1 <- testSet %>% mutate(default_payment =
                               factor(ifelse(default_payment == "1", "Yes", "No")))

# Scaling the Test Set for the Full Model

testSetFull <- predict(preProcess_range_model, newdata = testSet1)
testSet_under_sample1 <- predict(preProcess_range_model1, newdata = testSet1)

# Developing a MARS Model- Multivariate Adaptive Regression Spline to
# Predict the defaults on the Full data , in the later part
# we would be developing and Testing it on the undersample Data
# We would be developing this model using both Tune Length and Tune Grid on the
# Undersampled Training Set
# we would be Observing the Accuracy and the Sensitivity as we as the Specificity to
# understand the performance of the model
# Define the training control

fitControl <- trainControl(
  method = 'cv',                # k-fold cross validation
  number = 5,                   # number of folds
  savePredictions = 'final',    # saves predictions for optimal tuning parameter
  classProbs = TRUE,            # should class probabilities be returned
  summaryFunction=twoClassSummary # results summary function
)
```

```

# Step 1: Tune hyper parameters by setting tuneLength
set.seed(100)
model_mars1 <- train(default_payment ~ ., data=trainSet1, method='earth',
                     tuneLength = 5, metric='ROC', trControl = fitControl)
# model_mars1

# Step 2: Predict on testData and Compute the confusion matrix

predictedMars1 <- predict(model_mars1, testSetFull)

con_matrix_Mars <- confusionMatrix(predictedMars1, testSetFull$default_payment)

Mars.accuracy <- con_matrix_Mars$overall[["Accuracy"]]
Mars.specificity <- con_matrix_Mars$byClass[["Specificity"]]
Mars.sensitivity <- con_matrix_Mars$byClass[["Sensitivity"]]

# Doing the predictions Using a Tune Grid

# Step 1: Define the tuneGrid
marsGrid <- expand.grid(nprune = c(2, 4, 6, 8, 10),
                      degree = c(1, 2, 3))

# Step 2: Tune hyper parameters by setting tuneGrid
set.seed(100)
model_mars2 = train(default_payment ~ ., data=trainSet1, method='earth',
                    metric='ROC', tuneGrid = marsGrid, trControl = fitControl)
# model_mars2

# Step 3: Predict on testData and Compute the confusion matrix
predictedMars2 <- predict(model_mars2, testSetFull)

con_Matrix_Mars_Tuned <- confusionMatrix(predictedMars2, testSetFull$default_payment)

Mars_Tuned.accuracy <- con_Matrix_Mars_Tuned$overall[["Accuracy"]]
Mars_Tuned.specificity <- con_Matrix_Mars_Tuned$byClass[["Specificity"]]
Mars_Tuned.sensitivity <- con_Matrix_Mars_Tuned$byClass[["Sensitivity"]]

```

Developing the MARS model on the undersampled data

```

# Developing a MARS Model- Multivariate Adaptive Regression Spline to Predict the defaults
# on the under sampled Data
# We would be developing this model using both Tune Length and Tune Grid on the
# Undersampled Training Set
# we would be Observing the Accuracy and the Sensitivity as we as the Specificity to
# understand the performance of the model
# Define the training control

```

```

fitControl <- trainControl(
  method = 'cv',           # k-fold cross validation
  number = 5,             # number of folds
  savePredictions = 'final', # saves predictions for optimal tuning parameter
  classProbs = TRUE,      # should class probabilities be returned
  summaryFunction=twoClassSummary # results summary function
)

# Step 1: Tune hyper parameters by setting tuneLength
set.seed(100)
model_mars1_under_sample <- train(default_payment ~ ., data=train_under_sample1 ,
  method='earth', tuneLength = 5, metric='ROC', trControl = fitControl)
# model_mars1_under_sample

# Step 2: Predict on testData and Compute the confusion matrix

predictedMars1_under_sample <- predict(model_mars1_under_sample, testSet_under_sample1)

con_Matrix_Mars_Under_sample <- confusionMatrix(predictedMars1_under_sample,
  testSet_under_sample1$default_payment)

Mars_Under_Sample.accuracy <- con_Matrix_Mars_Under_sample$overall[["Accuracy"]]
Mars_Under_Sample.specificity <- con_Matrix_Mars_Under_sample$byClass[["Specificity"]]
Mars_Under_Sample.sensitivity <- con_Matrix_Mars_Under_sample$byClass[["Sensitivity"]]

# Doing the predictions Using a Tune Grid

# Step 1: Define the tuneGrid
marsGrid <- expand.grid(nprune = c(2, 4, 6, 8, 10),
  degree = c(1, 2, 3))

# Step 2: Tune hyper parameters by setting tuneGrid
set.seed(100)
model_mars2_under_sample = train(default_payment ~ ., data=train_under_sample1 ,
  method='earth', metric='ROC', tuneGrid = marsGrid, trControl = fitControl)
# model_mars2_under_sample

# Step 3: Predict on testData and Compute the confusion matrix
predictedMars2_under_sample <- predict(model_mars2_under_sample, testSet_under_sample1)

con_Matrix_Mars_Tuned_Under_sample <- confusionMatrix(predictedMars2_under_sample,
  testSet_under_sample1$default_payment)

Mars_Tuned_Under_Sample.accuracy <- con_Matrix_Mars_Tuned_Under_sample$overall[["Accuracy"]]

```

```
Mars_Tuned_Under_Sample.specificity <- con_Matrix_Mars_Tuned_Under_sample$byClass[["Specificity"]]
Mars_Tuned_Under_Sample.sensitivity <-con_Matrix_Mars_Tuned_Under_sample$byClass[["Sensitivity"]]
```

Creating the AdaBoost Model

```
set.seed(100)

# Train the model using adaboost
model_adaboost1 <- train(default_payment ~ ., data=trainSet1, method='adaboost',
                        tuneLength=2, trControl = fitControl)

# model_adaboost1

predicted_adaboost1 <- predict(model_adaboost1, testSetFull)
con_Matrix_adaboost <-confusionMatrix(predicted_adaboost1, testSetFull$default_payment)

Adaboost.accuracy <- con_Matrix_adaboost$overall[["Accuracy"]]
Adaboost.specificity <- con_Matrix_adaboost$byClass[["Specificity"]]
Adaboost.sensitivity <-con_Matrix_adaboost$byClass[["Sensitivity"]]
```

Creating and testing the adaboost Model on the Under Sampled Data

```
set.seed(100)

# Train the model using adaboost
model_adaboost1_under_sample = train(default_payment ~ ., data=train_under_sample1,
                                     method='adaboost', tuneLength=2, trControl = fitControl)

# model_adaboost1_under_sample

predicted_adaboost1_under_sample <- predict(model_adaboost1_under_sample, testSet_under_sample1)
confusionMatrix(predicted_adaboost1_under_sample, testSet_under_sample1$default_payment)
```

Confusion Matrix and Statistics

Reference

Prediction No Yes No 4288 586 Yes 1553 1073

Accuracy : 0.7148
 95% CI : (0.7044, 0.725)
 No Information Rate : 0.7788
 P-Value [Acc > NIR] : 1

Kappa : 0.3151

Mcnemar's Test P-Value : <2e-16

```

Sensitivity : 0.7341
Specificity : 0.6468
Pos Pred Value : 0.8798
Neg Pred Value : 0.4086
Prevalence : 0.7788
Detection Rate : 0.5717

```

```

Detection Prevalence : 0.6499
Balanced Accuracy : 0.6904

```

```
'Positive' Class : No
```

```

con_Matrix_adaboost_under_sample <- confusionMatrix(predicted_adaboost1_under_sample, testSet_under_sample)

Adaboost_under_sample.accuracy <- con_Matrix_adaboost_under_sample$overall[["Accuracy"]]
Adaboost_under_sample.specificity <- con_Matrix_adaboost_under_sample$byClass[["Specificity"]]
Adaboost_under_sample.sensitivity <- con_Matrix_adaboost_under_sample$byClass[["Sensitivity"]]

```

Compiling the Results of All the Models

Comparison of the Performance of the different Models in terms of the Overall Accuracy, Sensitivity and Specificity

```

The_Models <- c('Logistic_Regression_Cut_Off_0.5', 'Logistic_Regression_Cut_Off_0.25',
'Pruned_Tree_Under_Sample', 'Pruned_Tree_Prior_Probability',
'Pruned_Tree_Loss_Matrix', 'Pruned_Tree_Weights',
'MARS', 'MARS_Tuned',
'MARS_Under_Sample', 'MARS_Under_Sample',
'Adaboost', 'Adaboost_Under_Sample')

```

```

Accuracy_models <- c(log.accuracy.0.5 , log.accuracy.0.25 ,ptree_undersample.accuracy ,
ptree_prior.accuracy , ptree_loss_matrix.accuracy, ptree_weights.accuracy , Mars.accuracy ,
Mars_Tuned.accuracy, Mars_Under_Sample.accuracy, Mars_Tuned_Under_Sample.accuracy,
Adaboost.accuracy,Adaboost_under_sample.accuracy )

```

```

Specificity_models <- c(log.specificity.0.5 , log.specificity.0.25 ,ptree_undersample.specificity ,
ptree_prior.specificity , ptree_loss_matrix.specificity, ptree_weights.specificity , Mars.specificity ,
Mars_Tuned.specificity, Mars_Under_Sample.specificity, Mars_Tuned_Under_Sample.specificity,
Adaboost.specificity,Adaboost_under_sample.specificity )

```

```

Sensitivity_models <- c(log.sensitivity.0.5 , log.sensitivity.0.25 ,ptree_undersample.sensitivity ,
ptree_prior.sensitivity , ptree_loss_matrix.sensitivity, ptree_weights.sensitivity , Mars.sensitivity ,
Mars_Tuned.sensitivity, Mars_Under_Sample.sensitivity, Mars_Tuned_Under_Sample.sensitivity,
Adaboost.sensitivity,Adaboost_under_sample.sensitivity )

```

```

Results <- data.frame(The_Models, Accuracy_models,Sensitivity_models ,Specificity_models)

```


Giving The Summary of all the Modeld developed for easy comparison

```
kable(Results, col.names = c("The Models", "Accuracy ", " Sensitivity", " Specificity"),
      align = "cc", caption = "The Summary of All the Models Developed.")
```

Table 1: The Summary of All the Models Developed.

The Models	Accuracy	Sensitivity	Specificity
Logistic_Regression_Cut_Off_0.5	0.8065333	0.9684985	0.2362869
Logistic_Regression_Cut_Off_0.25	0.7556000	0.8144153	0.5485232
Pruned_Tree_Under_Sample	0.7390667	0.7639103	0.6515973
Pruned_Tree_Prior_Probability	0.7854667	0.8572162	0.5328511
Pruned_Tree_Loss_Matrix	0.4829333	0.3660332	0.8945148
Pruned_Tree_Weights	0.7776000	0.8353022	0.5744424
MARS	0.8146667	0.9453861	0.3544304
MARS_Tuned	0.8142667	0.9476117	0.3447860
MARS_Under_Sample	0.7584000	0.7979798	0.6190476
MARS_Under_Sample	0.7636000	0.8108201	0.5973478
Adaboost	0.8041333	0.9250128	0.3785413
Adaboost_Under_Sample	0.7148000	0.7341209	0.6467752

Conclusions

Here in the Models developed the PPositive is considered as Non Defaults Denoted by “0” and Negative as Default denoted by “1”.

We are here more Interested in the Defaults which are denoted by 1’s

Now in such a scenario as mentioned earlier the Sensitivity Denotes the Power of the Model , i.e the accuracy wrt the Predictions made, and Specificity denotes that out of the existing defaults , how accurately has the model been able to predict the defaults.

From the Banks point of view Specificity is more important here. In Such a Scenario the the Prune Tree Loss matrix and the Pruned Tree Under Sample gives better results Pruned Tree has the Best SPecificity but the accuracy is heavily compromised in order to achieve this.

The Under Sample Models of the Boosting Algorithms MARS and Adaboost giver better Specificity as compared to the Other Models while somewhat maintaining the Accuracies.

This is an unbalanced data set and these type of problems would be faced by the data Scientists designing the ML model to predict the Default.

The Reason for trying out so many model was to check all the options from a basic model to the Boosting algorithms, From Scaling the data to not scaling it and see what are the results whether there exists any significant difference between the output of a basic model to a very sophisticated model. There isnt any drastic difference , the difference is minor.

If the Power of the Model being a criteria where the option is that out of the predictions made , the predictions should be most accurate then MARS seems to be doing a Good Job. But if we want to identify the Maximum of the Defaulters than the models using an undersample data set for training does a decent job.

Here the company might have to define its cost , Of not predicting a would be defaulter correctly and what sort of default rate is sort of acceptable . Based on that Insight the cost component can be Included in

the model to give appropriate results which would be useful to the Organization which is the credit card company.

Based on the cost defined an appropriate model can be designed which would increase the relevant metric of interest to the organization , a Trade off might have to be decided upon.